

Artificial Intelligence

Albert-Ludwigs-Universität Freiburg



**UNI
FREIBURG**

Thorsten Schmidt

Abteilung für Mathematische Stochastik

www.stochastik.uni-freiburg.de

thorsten.schmidt@stochastik.uni-freiburg.de

SS 2017

Our goal today

Backpropagation

Regularization

Literature (incomplete, but growing):

- I. Goodfellow, Y. Bengio und A. Courville (2016). **Deep Learning**. <http://www.deeplearningbook.org>. MIT Press
- D. Barber (2012). **Bayesian Reasoning and Machine Learning**. Cambridge University Press
- R. S. Sutton und A. G. Barto (1998). **Reinforcement Learning : An Introduction**. MIT Press
- G. James u. a. (2014). **An Introduction to Statistical Learning: With Applications in R**. Springer Publishing Company, Incorporated. ISBN: 1461471370, 9781461471370
- T. Hastie, R. Tibshirani und J. Friedman (2009). **The Elements of Statistical Learning**. Springer Series in Statistics. Springer New York Inc. URL: <https://statweb.stanford.edu/~tibs/ElemStatLearn/>
- K. P. Murphy (2012). **Machine Learning: A Probabilistic Perspective**. MIT Press
- CRAN Task View: Machine Learning, available at <https://cran.r-project.org/web/views/MachineLearning.html>
- UCI ML Repository: <http://archive.ics.uci.edu/ml/> (371 datasets)

Forward propagation

- In a feedforward neural network to produce an output \hat{y} from an input x information flows forward through the network
- This is called forward propagation
- During training, forward propagation produces a scalar cost $J(\theta)$

Forward propagation algorithm for a typical deep neural net

- Require: Network depth, l
- Require: $W^{(i)}$, $i \in \{1, \dots, l\}$, the weight matrices of the model
- Require: $b^{(i)}$, $i \in \{1, \dots, l\}$, the bias parameters of the model
- Require: x , the input to process
- Require: y , the target output
- set $h^{(0)} = x$
- for $k = 1, \dots, l$ do:
 - $a^{(k)} = b^{(k)} + W^{(k)}h^{(k-1)}$
 - $h^{(k)} = f(a^{(k)})$
- at the end of the loop set:
- $\hat{y} = h^{(l)}$
- $J(\theta) = L(\hat{y}, y) + \lambda \Omega(\theta)$, where θ is $(W^{(i)}, b^{(i)})$ $i \in \{1, \dots, l\}$

- The back-propagation algorithm allows the information from the cost to flow backwards through the network, in order to compute the gradient
- The term back-propagation is not the whole learning algorithm
- Back-propagation is only a method to compute the gradient
- Another algorithm, e.g. stochastic gradient descent, is used to perform learning using this gradient.

- Computing an analytical expression for the gradient is straightforward
- Numerically evaluating such an expression can be computationally expensive
- The back-propagation algorithm does so using a simple and inexpensive procedure, that relates to the chain rule.

$$\begin{aligned} & \frac{\partial z}{\partial w} \\ &= \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial w} \\ &= f'(y) f'(x) f'(w) \\ &= f'(f(f(w))) f'(f(w)) f'(w) \end{aligned}$$




Figure from Goodfellow 2016

Backward propagation algorithm for a typical deep neural net

After the forward computation, compute the gradient on the output layer:

$$\mathbf{g} \leftarrow \nabla_{\hat{\mathbf{y}}} J = \nabla_{\hat{\mathbf{y}}} L(\hat{\mathbf{y}}, y)$$

for $k = l, l - 1, \dots, 1$ **do**

Convert the gradient on the layer's output into a gradient into the pre-nonlinearity activation (element-wise multiplication if f is element-wise):

$$\mathbf{g} \leftarrow \nabla_{\mathbf{a}^{(k)}} J = \mathbf{g} \odot f'(\mathbf{a}^{(k)})$$

Compute gradients on weights and biases (including the regularization term, where needed):

$$\nabla_{\mathbf{b}^{(k)}} J = \mathbf{g} + \lambda \nabla_{\mathbf{b}^{(k)}} \Omega(\theta)$$

$$\nabla_{\mathbf{W}^{(k)}} J = \mathbf{g} \mathbf{h}^{(k-1)\top} + \lambda \nabla_{\mathbf{W}^{(k)}} \Omega(\theta)$$

Propagate the gradients w.r.t. the next lower-level hidden layer's activations:

$$\mathbf{g} \leftarrow \nabla_{\mathbf{h}^{(k-1)}} J = \mathbf{W}^{(k)\top} \mathbf{g}$$

end for

this is Algorithm 6.4 in Goodfellow 2016

- The gradients on weights and biases can be used for a stochastic gradient update
- Symbol-to-number differentiation (Torch, Caffe): Use a set of numerical values for the inputs and return a set of numerical values describing the gradient at those input values
- Symbol-to-symbol differentiation (Theano, Tensorflow): Add additional nodes to the graph that provide a symbolic description of the desired derivatives.
- Because the derivatives are just another computational graph, it is possible to run back-propagation again, to obtain higher derivatives.




Figure from Goodfellow 2016

- regularization is a way to overcome underfitting, overfitting issues by trading variance of the prediction error against bias.
- $\mathbb{E}[L(\hat{y}, y)] = \text{Irreducible Error} + \text{Bias}^2 + \text{Variance}$ (exercise)
- regularization is a modification to a learning algorithm that is intended to reduce its generalization error but not its training error.
- we have already seen bagging as a regularization method
- In the context of deep learning, most regularization strategies are based on regularizing estimators, by adding a parameter norm penalty $\Omega(\theta)$ to J

$$J(\theta; X, y) + \lambda \Omega(\theta)$$

- weight decay refers to the L^2 penalty.
- also known as ridge regression
- if we do not punish the bias b the objective function for weight decay is given by

$$\tilde{J}(w; X, y) = \frac{\lambda}{2} w^T w + J(w; X, y)$$

- this means in a single gradient update step the update changes to

$$w \leftarrow (1 - \varepsilon \lambda) w - \varepsilon \nabla_w J(w; X, y)$$

- the addition of the weight decay term has modified the learning rule to shrink the weight vector on each step

- we make a quadratic approximation to the objective function in the neighborhood of the value w^* , the optimal weights where unregularized training cost is minimal

$$\hat{J}(w) = J(w^*) + \frac{1}{2}(w - w^*)^T H(w - w^*)$$

- where H is the Hessian matrix of J with respect to w evaluated at w^*
- the minimum of the regularized version of \tilde{J} is at

$$\tilde{w} = (H + \lambda I)^{-1} H w^*$$

- If we decompose $H = Q\Lambda Q^T$ into a diagonal matrix Λ and an orthonormal basis of eigenvectors Q we get

$$\tilde{w} = Q(\Lambda + \lambda I)^{-1} \Lambda Q^T w^*$$




Figure from Goodfellow 2016

- In comparison to L^2 regularization, L^1 regularization results in a solution that is more sparse.
- Sparsity in this context refers to the fact that some weights have an optimal value of zero.

Let us have a look at the learning procedure at `playground.tensorflow.org`